

Agility in the Age of AI

Sarah McConnell | MML (Measure Model Lead)

The Pattern We've Seen Before

Agile began as a genuine response to a real problem. Small teams, close to the work, figured out that iterative delivery in short cycles produced better software than long upfront plans. It worked for them. Others noticed, borrowed the ideas, and over time a methodology emerged. Then the bandwagon started. Consultancies packaged it. Certifications appeared. Executives heard it would make their organisations faster and cheaper to run. And the scramble to adopt it began - not because organisations had deeply understood the problem it was solving, but because everyone else was doing it and the pressure to keep up was real.

Most organisations did the same thing. They restructured delivery teams, introduced sprints and standups, sent people on courses, and reported upward that the transformation was underway. What they did not change was everything above the team level. Leadership continued to operate in annual planning cycles. Decisions stayed at the top. The impediments agile surfaced - unclear priorities, unresolved dependencies, years of accumulated technical debt - were logged in retrospectives and left unresolved, because resolving them required authority that teams did not have and leaders were not prepared to delegate.

Scrum at the bottom. Waterfall at the top. The vocabulary of transformation without the substance of it.

It happened for understandable reasons. Genuine agile transformation is slow, uncomfortable and politically difficult. Buying a framework and running some ceremonies is faster and easier to report upward. Revenue pressure and competitive pressure make the surface version attractive. So most organisations took it.

The data reflects this. Standish Group data shows roughly 58% of agile transformations either failed outright or were significantly challenged. McKinsey found only 4% of companies had completed a genuine enterprise-wide agile transformation by 2017, a decade into mainstream adoption. Digital.ai's 17th State of Agile Report found only 11% of practitioners highly satisfied with their organisation's implementation, with lack of leadership participation cited as the single biggest barrier. These are not failure rates caused by bad frameworks. They are failure rates caused by organisations that applied the surface of a transformation without doing the systemic work underneath it.

The organisations that got it right did something harder. They treated agile as a whole-system question - which meant changing how decisions were made and at what level, investing in removing the blockers their teams raised rather than absorbing them, ensuring customer insight actually reached the people building the product, and managing technical debt as an ongoing discipline rather than a problem to defer. They brought in people whose job was not to implement a framework but to understand their specific organisation and help it find its own way through. The framework was a starting point. The transformation was the work.

Most organisations did not do that work. And now they are doing the same thing again, with AI.

AI followed the same path, but faster and louder. Machine learning produced genuine breakthroughs in narrow domains. Technology companies scaled the tooling. And then the narrative shifted in a way agile never quite managed: from "this could give you an advantage" to "not doing this will make you obsolete." McKinsey's State of AI 2025 states it plainly - AI is no longer a competitive advantage, it is the baseline. That framing created a different quality of pressure to anything agile produced. Boards began asking questions. Investors began asking questions. And the fastest way to demonstrate an answer was to procure a tool and push staff to use it.

A tool arrived. Comms went out. Staff were encouraged to use it. Adoption metrics were tracked and reported upward. And six months later, when someone asked what had actually changed in how the organisation delivered, the honest answer was quietly uncomfortable. The tool was there. The transformation was not.

MIT's Project NANDA reviewed over 300 AI initiatives in 2025 and found 95% of generative AI pilots delivered zero measurable P&L impact. RAND puts the AI project failure rate at over 80%. BCG found 74% of organisations struggling to achieve and scale AI value. IBM's 2025 CEO study found only 25% of AI initiatives had delivered expected ROI.

These are not technology failures. The tools work. What is failing is the organisational substrate AI is being asked to land on. It is the same substrate agile was asking organisations to fix. Most of them never did.

What AI Transformation Actually Requires

Organisations that are serious about AI transformation will find themselves facing a familiar set of problems. The good news is these are not new. The bad news is that many companies failed to answer them last time round - the same ones agile surfaced and most organisations chose not to fully resolve.

The difference is that AI makes the cost of leaving them unresolved significantly higher, and significantly more visible.

The technical foundations

AI cannot optimise a broken system. It can only move through it faster, which means broken processes produce broken outputs at greater speed and scale. Two infrastructure problems sit underneath almost every AI failure that gets attributed to the technology itself.

The first is technical debt. Accumulated shortcuts, legacy systems, undocumented dependencies, deployment pipelines held together with workarounds - these are not AI problems. They predate AI by years or decades in most organisations. But AI lands on them and fails. IBM's 2025 research found 86% of executives say technical debt is actively constraining their AI success. Google engineers identified as far back as 2015 that machine learning systems inherit all of the maintenance problems of traditional code plus an entirely new set of ML-specific ones. The organisations

treating AI as a reason to finally address their technical debt are the ones making genuine progress. The ones treating it as something to work around are building on quicksand - organisations that find themselves here will only sink faster if they continue to thrash about rather than stop, think and find the right solution.

The second is deployment pipelines. Most organisations have not solved continuous delivery even for traditional software. DORA's 2024 State of DevOps report found only 19% of teams reach elite performance levels, and the proportion of high performers actually shrank year on year. Adding AI into an immature pipeline does not accelerate delivery. It adds more traffic to an existing bottleneck. The same DORA research found AI coding tool adoption correlated with a drop in throughput and stability in teams with weak delivery foundations - the opposite of what was promised.

These are not glamorous problems. They do not make for exciting board presentations. But they are the foundations everything else depends on, and organisations that skip them will keep producing the same outcome - impressive pilots, negligible enterprise impact.

People, organisation and culture

BCG's research across 1,000 executives established that roughly 70% of AI transformation challenges are people and process problems, 20% are tools and process, and only 10% involve the actual algorithms. That figure surprises some people. It shouldn't. Consider what sits inside that 70%: how decisions are made, whether teams have genuine authority, whether customer insight reaches the people who need it, whether leaders enable or obstruct, and whether the organisation has created the conditions for people to work honestly with new systems rather than around them.

Decision-making is where AI transformations stall most visibly. McKinsey's research found only 37% of organisations make decisions that are both high quality and high velocity. For delegated decisions the figure drops to 27%. AI amplifies this problem because it operates at a speed that exposes slow, hierarchical decision-making immediately. An organisation where meaningful calls travel up a chain and come back too slowly was already losing ground. In an AI-augmented system it loses ground faster. Genuine empowerment is not a cultural nice-to-have. It is what makes fast, good decisions possible at scale. McKinsey found that organisations where employees are genuinely empowered are three times more likely to achieve both. That is not a technology investment. It is a leadership one.

Customer insight follows the same pattern. Teresa Torres, whose work on continuous product discovery has become a reference point for modern delivery organisations, found that product teams are often four degrees removed from their customers - with sales, marketing and research functions acting as intermediaries who filter, delay and shape what reaches the people building the product. AI models trained on siloed, filtered or biased data produce siloed, filtered and biased outputs. MIT Sloan's research identifies data accessibility as primarily an organisational problem, not a technical one. The data exists in most organisations. The structure prevents it from flowing to where it is needed.

Leadership behaviour is the multiplier across all of it. Poor leadership in an AI transformation is not subtle. It looks like grand announcements made in isolation, landing on people who are already worried about their jobs with no honest conversation about what is actually changing or why. It looks like pilot operating models designed by small groups and presented to the organisation as fait accompli, with the diversity of perspective that might have stress-tested them nowhere in the room. And increasingly it looks like AI-generated strategy presented as human thinking - frameworks and role definitions that carry the confident surface of considered work without the robustness of having actually discussed and defended it in a room. It looks like covert headcount planning that everyone suspects and nobody will confirm, while the leadership team asks those same people being replaced to embrace the change. That last point also carries legal risk most have not considered. The Employment Rights Act 2025, effective April 2026, doubled the maximum protective award for collective redundancy consultation failures to 180 days' full pay per affected employee, uncapped. Organisations making headcount decisions on the back of AI adoption without proper process are operating in a significantly higher-risk legal environment than they were twelve months ago, whether they know it or not.

The consequence is not just morale damage, though that is real. It is that the people who understand how the organisation actually works - where the real dependencies sit, what the customers actually need, why the last three initiatives did not deliver - are disengaged from the design of something that depends entirely on their knowledge and cooperation to succeed. BCG's finding that AI leaders combine strong leadership with decentralised execution is not an aspiration. It is a description of the only conditions under which this works.

Which brings culture into its proper place - not as a separate consideration but as the result of everything above done well or badly. Teams with psychological safety flag when an AI output is wrong. Teams without it produce it anyway and say nothing. Organisations that have been honest with their people about what is changing find those people working with new systems. Organisations that have not find them working around them. Harvard Business Review's 2025 research found that fear of replacement, rigid workflows and entrenched power structures are the primary organisational barriers to AI adoption - none of which show up on a technology roadmap, and all of which are cultural conditions that leadership either creates or allows.

When BCG says 70% of the challenge is people and process, they are describing an organisation's capacity to change how it works at a human level. That is not a softer challenge than fixing the technical foundations. It is a different kind of hard. And it is the kind of hard that cannot be outsourced to a tool.

What AI requires that is genuinely new

AI is simultaneously breathtakingly capable and unreasonably stupid. Part of that stupidity is particularly dangerous: it will complete the task you gave it without flagging that the foundations were shaky, the data was compromised, or the answer it produced has a fundamental flaw. There is no moment where it pauses and says "you may want to check this." It just delivers.

A Monte Carlo forecast is a methodology that runs thousands of simulations across a range of variables to model probable outcomes - and anyone who has built one in a spreadsheet will understand immediately why it becomes an AI candidate. Connect your AI to your data source, run the simulation, get a result. What the manual version forced you to do was enter that data yourself - and in doing so, see immediately that it was inconsistent, incomplete, or unreliable. The AI skips that moment entirely. It finds a way to produce a number regardless. The mess that would have stopped a human cold becomes invisible, and the answer arrives looking clean.

What this means in practice is that AI democratises access to sophisticated methodologies that previously required specialist knowledge to run. That sounds like a benefit. And it is, but it comes with a risk.

The risk is that the people who understood the methodology well enough to know when the output was wrong are no longer the ones running it. It is the leader four steps removed from the work who connected the tool, got a number, and is now telling the team they have ten weeks to deliver something that anyone close to the data would know will take four months. The forecast looks authoritative. Nobody in the room has the context to challenge it. And the organisation plans around a number that was wrong before it was even finished generating.

The democratisation of capability that AI enables is genuinely powerful - but only in organisations that have done the foundational work. Without data quality, governance and the human expertise to validate outputs, it does not democratise capability. It democratises confident wrongness at scale.

The difference between an organisation that catches this and one that doesn't is not the AI. It is whether the human reviewing the output understands enough about the methodology, the data sources and the known limitations to open the black box rather than accept what came out of it. That is not a technical skill. It is a judgment skill, a domain knowledge skill, and an organisational culture question about whether it is safe to challenge an output that leadership has already celebrated.

The feeling of involvement is not the same as being close to the work. Leaders and teams who believe they are directing AI because they craft the prompts are often further from the actual decision-making than they realise. The work is happening in a black box. The human feels present. The human is not fully in control. And in organisations where nobody is asking how they would know if the system was wrong, that gap between feeling in control and being in control is where the most expensive mistakes will be made.

This requires organisations to answer questions they have not had to answer before. What can AI be trusted to decide, and to what degree? Where does human judgment remain non-negotiable? Who owns accountability for an outcome when an agent produced it? What are the boundaries of AI autonomy within your specific organisational context, and how are those boundaries enforced?

These are governance questions, risk questions, and leadership questions. And they require genuine investment - proper frameworks, proper oversight, proper accountability structures - not because a regulator requires it, but because without it the system is ungoverned and the people working inside it are exposed.

Risk frameworks for AI are genuinely hard. Not because of their complexity alone, but because they require organisations to make explicit decisions about trust, autonomy and accountability that most have never had to make before. The organisations doing this well are treating it as a discipline in its own right, not a box to tick before the pilot goes live.

The operating model

All of the above - technical foundations, decision-making, governance, leadership, culture - cannot be addressed in isolation. They are not a checklist. They are components of an operating model, and AI transformation ultimately requires organisations to redesign how work is structured, not just how it is done. That distinction matters because most organisations treat these as separate workstreams: a technical debt programme here, a culture initiative there, a governance framework bolted on at the end. Run in isolation, each produces local improvement and negligible enterprise impact. The failure mode is not that organisations ignore the problems. It is that they address them one at a time, in the wrong order, without ever asking how they connect.

The organisations that will realise genuine value from AI are not the ones that adopted it fastest. They are the ones that used AI as the forcing function to finally do the systemic work they should have done when agile first asked them to do it. Most organisations know what that work is. They have known for some time. The question AI is now forcing - more urgently than agile ever did - is whether this is the moment they decide to actually do it.

Where the Parallel Has Limits

The argument made in this paper deserves to be stress-tested, including by the person making it. There are two places where the parallel genuinely breaks down, and intellectual honesty requires naming them directly.

The first is pace. Agile was designed for human-speed delivery. The practices it produced - sprints, standups, retrospectives, reviews - exist because humans need rhythm, shared context and regular moments to check they're still pointing in the same direction. Agentic AI operates at machine pace. An agent can produce in minutes what a team previously produced in a sprint. The governance challenge that creates is not a scaled-up version of the agile governance challenge. It is a different kind of problem entirely. How do you maintain meaningful human oversight of a system that moves faster than any review cadence designed for humans? How do you define done for a system whose outputs shift as the underlying data shifts? Agile does not answer those questions. Nobody has fully answered them yet.

The second is the vendor data. MIT's 2025 research found vendor-partnered AI succeeds at roughly twice the rate of internally built AI. That gap is real. Vendors typically succeed because they are solving a specific, contained problem with deep specialist expertise - and they are genuinely good at it. That is a different activity to transforming how an organisation operates at scale, which is what this paper is about. Both are legitimate. They are not the same problem.

These are genuine limits. The parallel does not cover everything. And the more considered version of the argument - the one that survives scrutiny - is not that agile and AI transformation are the same thing. It is that the organisational conditions agile requires are a prerequisite for AI transformation, and most organisations have not met them. What comes after those prerequisites are met is harder, newer, and less well understood. That is where the real work is.

What This Means In Practice

In an ideal world, the first question any organisation asks before committing to AI transformation is not "what is our AI strategy" but "do we understand our own system well enough to know where AI belongs in it." The two questions can feel like the same question. They are not. A strategy sets direction. The question that should precede it is investigative: it requires getting close enough to the work to know where value is actually created and lost, where the real bottlenecks sit, which decisions are genuinely high-leverage and which only feel that way, and where human judgment is non-negotiable versus where it is being applied to work that does not need it. A strategy built without that understanding is not wrong because the thinking is poor. It is wrong because it inherently carries risk. The risk of a shaky foundation that, when optimised for AI, swiftly carries a business in the wrong direction.

Understanding where you stand starts with understanding your organisation holistically. Where are decisions actually made, and how fast do they move? Can the people building products get to the data they need without filing a request and waiting three weeks? Is technical debt being managed as a discipline or quietly deferred? Are leaders creating the conditions for honest work, or generating the kind of low-grade distrust that causes people to work around new systems rather than with them? These are not new questions. Anyone who has been through a serious agile transformation will recognise every one of them. That is not a coincidence. The signal you are looking for in an AI transformation is the same signal, in the same places, as the one you have seen before. For organisations already mid-journey, that signal is still readable - in the tools that aren't delivering, the pilots that failed quietly, the adoption that never stuck. Failure is information. It points to the same places these questions point to.

Organisations that have achieved genuine business agility - not the ceremony layer, but the substance of it: distributed authority, short feedback loops, customer insight that actually reaches the people who need it, technical debt managed rather than deferred - have done most of this work already. They will not find AI transformation easy, but they will find it legible. The ones that ran the ceremonies without the underlying redesign are facing the same ask they declined a decade ago. The stakes are higher. The window is shorter. But the work is the same work, and knowing that is more useful than pretending it is something entirely new.

The same principles apply with more force when organisations move from AI tools to AI agents - systems that do not just assist decisions but make them, autonomously, at machine pace. This is where the organisational design question becomes critical rather than just important. Consider how a well-functioning self-governing team operates: clear boundaries defining what sits within its authority and what does not, genuine autonomy inside those boundaries, and accountability for outcomes rather than sign-off on activity. The boundaries are not constraints on what the team can

achieve - they are what makes it safe to give the team real authority in the first place. Holacracy, the governance model that formalised these principles into a full organisational design, demonstrated that distributed authority works when the boundaries are explicit and the accountability is real. Agentic AI requires exactly the same design logic. The guardrails - what the agent can do, what it cannot, under what conditions it escalates - are not limitations on capability. They are the architecture that makes genuine autonomy possible without the organisation losing the ability to trust what the system is doing. An organisation that has never built that architecture for its human teams will not build it for its AI agents. It will either over-constrain the agents until they are useless, or under-constrain them until something goes wrong.

The organisations that make genuine progress share a common characteristic in their leadership: people who treat AI evolution as something to lead rather than manage. That distinction matters more here than in almost any other context, because the gap between the two is where most AI transformation value is lost. Managing AI - controlling outputs, approving decisions, maintaining oversight at the activity level - creates exactly the bottleneck that makes AI underperform. The machine moves at machine pace. The approval chain does not.

Leading AI evolution means something more specific than delegation. It requires a combination of strong leadership, decentralised execution and shared ownership between business and technology. BCG's research across 1,400 C-suite leaders found only 5% of companies are capable of doing this today. And this is not a new concept - Marty Cagan published "Empowered" in 2020 and states "moving to empowered teams does not mean fewer leaders, it means stronger ones." Strong leaders are able to set direction clearly enough that good decisions can be made without them, they hold people accountable for outcomes rather than activity, and they design the conditions under which the organisation can move at the speed the technology makes possible. Successful AI implementation needs strong leadership, not just people who are making technology choices. The organisations whose leaders can meet this challenge are the ones with something to show for their investment.

Agile has asked this of leaders for decades: get close to the customer, push decisions down to where the work happens, invest in the foundations rather than deferring them, and make the feedback loop short enough to learn from it. AI is asking for all of that, and one thing more: decide, explicitly and in advance, what should never be delegated to a machine at all.

Most organisations did some of this. Very few changed their foundations. AI has considerably less tolerance for a skim coat of change, and finding that out is expensive. The ones that recognise that will come out the other side with something that actually works. Those that don't will have very impressive pilots.

Appendix: The Harder Questions: Organisational Conditions for AI Transformation

The paper has made the case that AI transformation fails for organisational reasons, not technological ones. The data is consistent across agile and AI: high failure rates, low enterprise impact, and a gap between what organisations report and what they

deliver. The root causes are not new. They are structural, behavioural and organisational - and most of them were visible long before the technology arrived.

What is harder is knowing where those root causes sit in your specific organisation. The failure patterns described in this paper are common. The way they manifest varies. An organisation with strong delivery maturity but weak decision-making will hit different walls than one with empowered teams and a data accessibility problem. Understanding which conditions are present, and how severely, is the work that precedes everything else.

The domains that follow are the places where the gap between where an organisation thinks it is and where it actually is tends to be widest. They are not exhaustive. Each one covers why it matters, how weakness tends to show up in practice, and a set of questions worth asking honestly. Some of those questions will be uncomfortable.

This is not a substitute for doing the work. It is a starting point for understanding what the work actually is.

Decision-making and delegated authority

Why it matters

AI operates at a pace that makes slow, hierarchical decision-making both inefficient and structurally incompatible with what is being asked of it. Every organisation has a decision-making process. What varies - and what matters - is the speed at which decisions are made and the level at which they are made. Both have direct bearing on which processes are viable candidates for AI involvement. Add to that the developing governance landscape around what AI should and should not be permitted to do, and organisations that have not thought this through are exposed to operational and legal risk they have not accounted for.

How weakness tends to show up

The most common sign is a persistent gap between what leadership believes is possible and what teams can actually deliver. Strategy gets agreed and then quietly stalls somewhere between the room it was decided in and the people expected to implement it. Planning cycles repeat the same conversations without resolution. The cause is rarely visible on a dashboard - it shows up in low team morale that nobody can quite explain, and in the fact that when you ask someone to describe how decisions actually flow in their organisation, most cannot do it with any confidence.

Questions worth asking

Where are decisions actually made, and where are they supposed to be made?

How long does it take for a decision that needs to be made at team level to be made and acted on?

When decisions are delegated, is the authority real or conditional?

What happens when someone makes a delegated decision that turns out to be wrong?

Technical debt and delivery maturity

Why it matters

AI promises faster, leaner development of code and products. What it cannot do is fix a broken deployment system or shaky technical foundations. What it will do is surface them, quickly and visibly. Significant technical debt, immature pipelines and undocumented dependencies do not disappear when AI arrives. They become the first thing it runs into. That is not a failure of the technology. It is the accumulated consequence of a stack that has been neglected, deferred and worked around for long enough that the workarounds have become the system.

How weakness tends to show up

Lead times are high and have been for long enough that nobody questions them anymore. DevOps initiatives have been attempted and have not stuck. CI/CD exists in parts of the system but not consistently across it. Technical debt has been chronically underfunded, and the changes required to genuinely address it are not quick fixes - some carry timelines measured in years, not months. Testing and reference environments are out of date compared to production. AI pilots work in controlled conditions and fail when they meet the full production system.

Questions worth asking

How long does it take to deploy a change from code complete to production?

What proportion of engineering time goes on maintenance and workarounds versus new capability?

What is on the technical debt backlog that has been there for more than two years, and why is it still there?

Does a current, accurate map of your system dependencies exist, and when was it last validated?

Access to data and customer insight

Why it matters

Despite the name, AI is not yet intelligent in any meaningful sense. The models organisations are building on are fundamentally dependent on the quality of what goes into them. Poor data, filtered insight and inaccessible customer reality produce outputs that accelerate confidently in the wrong direction. An intentional strategy around data quality and customer insight is what separates AI that deepens your understanding of your customers from AI that hardens your misconceptions about them. Your AI will learn from the data it is served, and without intervention this problem does not stay static. It grows.

How weakness tends to show up

Data exists across most organisations in abundance - marketing holds survey responses, support owns ticket data, product tracks feature requests - but ownership is fragmented across functions that have no incentive to share it and no mechanism to do so. The research agrees: MuleSoft's 2025 Connectivity Benchmark found that only 2% of businesses have integrated more than half their applications. No single person has a complete picture of what customers actually experience. Gartner estimates that over 87% of organisations struggle with disconnected data sources.

When AI is trained on one of these fragmented sources it produces outputs that reflect that function's view of the customer, not the customer's actual experience.

The operational consequence is that AI fails not because the model is wrong but because what it was given to learn from was incomplete. Teams build on outputs they cannot validate because they do not have access to the data that would allow them to challenge them. The insight exists somewhere in the organisation. The structure ensures it never reaches the people who need it.

Questions worth asking

How many steps exist between your customers and the people building for them?

When did someone on your product or delivery team last speak directly with a customer?

If a team needed a specific dataset today, how long would it take to get it and who would they ask?

Are there data sources your organisation holds that the people who would most benefit from them cannot reach?

Leadership behaviour, trust and psychological safety

Why it matters

The media narrative around AI is not subtle. Predictions about significant job losses are everywhere, and employees are paying attention. Unlike previous transformation programmes, AI is genuinely new ground. No organisation has a complete map for what it looks like at scale, which means leaders are responsible for communicating honestly about a journey whose destination they cannot yet fully describe. The ones who do create the conditions for people to engage with the uncertainty. Leaders who do not leave the narrative to fill the gap for them. The consequence is not just low morale. It is employees who do not buy into AI initiatives, or who actively work against them.

How weakness tends to show up

AI transformation is treated like any other major organisational change - decisions made by a small group, implementation handed downward. The difference is that unlike most change programmes, it arrives with a broadcast. Announcements, enthusiasm, external communications about innovation and ambition. The gap between that signal and the absence of any honest conversation about what it means for the people receiving it is where trust breaks down. Questions about roles and headcount go unanswered or are deflected. Without a visible roadmap and genuine transparency, the narrative leaves leadership's hands entirely. Low morale is the early indicator. Over time the gap between what leadership believes about employee sentiment and what employees actually feel becomes wide enough that initiatives start failing for reasons nobody is officially naming.

The signs are rarely dramatic and rarely named. The organisation looks like it is functioning. Metrics are being hit, tools are being used, retrospectives are happening. But the quality of honesty in the room has quietly degraded. People have learned what is safe to say and are saying that. The voices that go quietest are

usually the ones with the most to offer - the people who understand why things fail and have stopped volunteering that understanding because the environment no longer makes it worth their while. Some of them leave. The gap between the official version of how the AI transformation is going and the version that exists in informal conversations grows steadily, and nobody with the authority to close it is aware of how wide it has become.

Questions worth asking

When someone raises a concern about an AI output or process, what actually happens?

How were the people most affected by your AI initiatives involved in designing them?

Is there an honest, shared understanding in your organisation of what AI adoption means for roles?

Where do you see evidence that people are going through the motions with AI tools rather than genuinely using them?

Governance, accountability and AI risk boundaries

Why it matters

The regulatory environment is moving faster than most organisations' governance frameworks. The Data (Use and Access) Act 2025 introduced mandatory safeguards for automated decision-making in the UK - requiring organisations to inform individuals of significant decisions made about them, enable challenges to those decisions, and provide meaningful human intervention. The ICO's guidance on what constitutes meaningful human involvement is in public consultation now, with final guidance expected Summer 2026.

For organisations with EU employees or customers, EU AI Act obligations covering high-risk AI systems come into force August 2026, adding a parallel compliance layer that the UK framework does not automatically satisfy. Organisations that have not mapped where AI is making or influencing significant decisions about individuals are already operating in regulated territory without the infrastructure to demonstrate compliance.

Organisations deploying AI are also taking on accountability they may not have fully considered. In February 2024, a BC Civil Resolution Tribunal ruling in *Moffatt v. Air Canada* established the first major Commonwealth legal precedent for corporate liability for AI outputs, rejecting the argument that an AI system could be treated as a separate legal entity responsible for its own actions. The organisation was liable.

Gartner's April 2026 research found that over 40% of agentic AI projects will be cancelled by end of 2027, with insufficient governance infrastructure and absent accountability structures among the leading reasons. Organisations that think through these questions before deployment protect themselves from liability and give their AI initiatives a significantly better chance of success.

How weakness tends to show up

More attention goes on what to do with AI than on how to integrate it safely. AI gets treated as another participant in an existing system rather than a fundamentally different way of working that requires its own risk and governance layer - one for which frameworks are still in early development. The most telling sign is the question that does not get asked: not what AI can do, but what it should do. The boundaries of what AI is permitted to do in a given context are assumed rather than defined, which means they are not boundaries at all.

Questions worth asking

If an AI system in your organisation produced a consequential error today, who would be accountable and through what process?

Where is AI currently making or materially influencing decisions, and is that list complete and current?

How would you know if an AI system was producing systematically biased or incorrect outputs?

What is the process for defining what AI is and is not permitted to do in a given context, and who owns it?

What should remain human judgment

Why it matters

AI has fundamental gaps that are revealed when it is allowed to own the end to end of a process which requires decisions. It does not know what it does not know. But it will produce a confident, coherent answer by default and proceed with its task. The decisions that require contextual nuance, domain knowledge, and the judgment to ask whether the correct answer is actually the right one are not ones AI can make safely. Without a human who is genuinely accountable for the outcome, an organisation is exposed. Most organisations have not identified which decisions are suitable, or defined a suitability criteria for agentification or augmentation of processes. The question "what should AI own, and what should it never own" sits underneath all of the other work in this appendix. Without an answer to it, the guardrails don't exist to protect the organisation.

In 2020, the Ofqual grading algorithm standardised A-level grades against historical school performance - exactly what it was designed to do. Any human would question whether it is fair to use past performance as a measure for students being educated during a global pandemic. The algorithm had no mechanism to come up with that question. Around 40% of students received lower grades than their teachers predicted, disproportionately affecting those at large state schools. The government reversed course under public pressure before legal challenges reached court.

This kind of U-turn is not available to a commercial organisation that has already hired, fired, lent, denied or restructured on the basis of a flawed output. Having a human in the process who would have known to ask the question first, or who is accountable for skipping assurance steps, protects the organisation by reducing the risk of these situations arising at all.

How weakness tends to show up

The clearest sign is the absence of a problem statement. AI is adopted as a standard of working rather than as a solution to a specific, defined need - which means nobody has asked which problems it is actually solving or whether those problems required AI at all. Governance and oversight are treated as a later consideration rather than a precondition. The boundary between what AI should assist and what it should determine is never drawn explicitly, which means it gets drawn by default, usually by whoever is moving fastest.

The longer term sign is what leaves when people do. Redundancies made on the assumption that AI can absorb the work remove domain knowledge, process understanding and stakeholder relationships built over years. The institutional memory that would have caught a bad output or known which edge cases needed different handling goes with the people who held it. Organisations rarely notice the gap until they need to question something the system is telling them and find they no longer have anyone who can.

Questions worth asking

Which decisions in your organisation require human judgment, and does everyone involved know which ones those are?

Where is AI currently influencing decisions in ways that would surprise the people affected by them?

What would need to be true for you to be confident that a human reviewing an AI output was genuinely reviewing it rather than ratifying it?

Are there use cases in development or deployment where the consequences of a systematic error have not been fully examined?

Research dossier available separately.

Sarah McConnell | MML (Measure Model Lead)